

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import statsmodels.api as sm
import statsmodels.formula.api as smf

# Load the dataset
insurance_df = pd.read_csv("/Users/balakrishnamupparaju/Downloads/insurance.csv")

# Inspect the dataset
print(insurance_df.info())
print(insurance_df.head())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
None
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
In [3]: # Check for missing values
print(insurance_df.isnull().sum())

# Data types of variables
print(insurance_df.dtypes)
```

```
age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
dtype: int64
age          int64
sex          object
bmi          float64
children     int64
smoker       object
region       object
charges      float64
dtype: object
```

```
In [5]: print(insurance_df.columns)
```

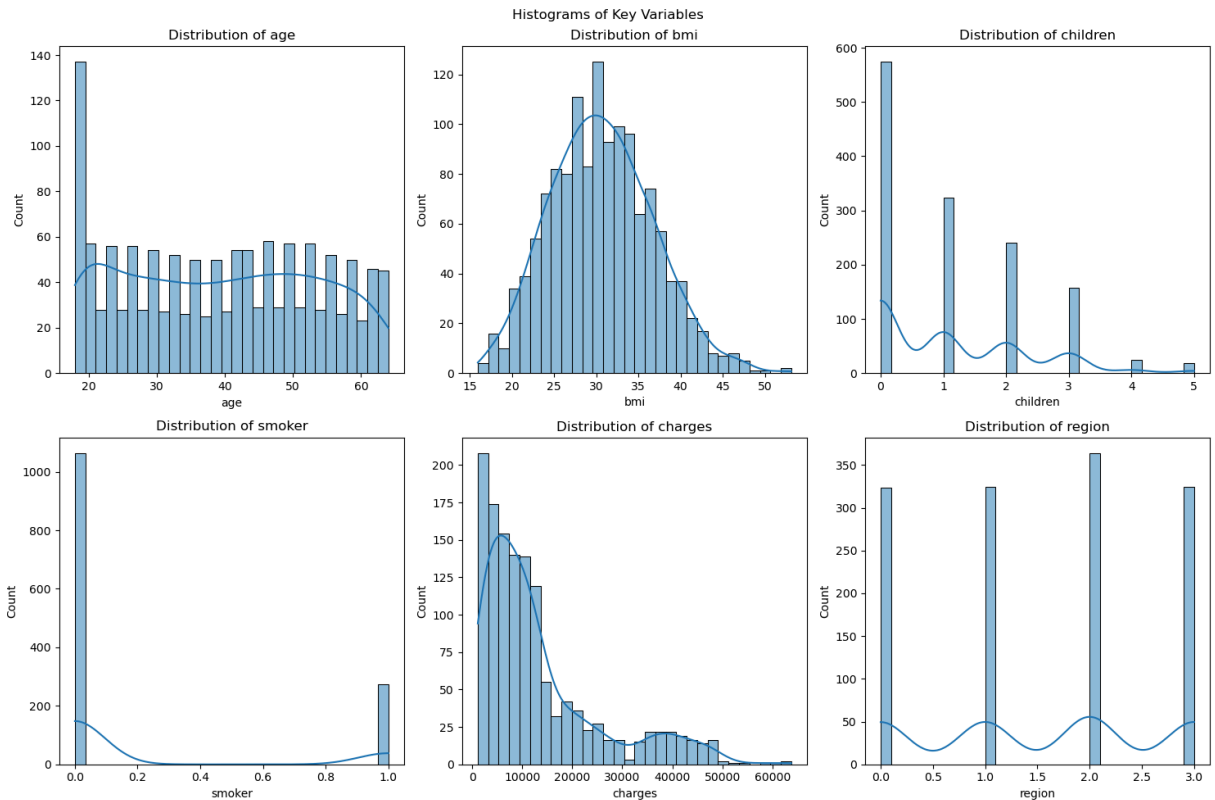
```
Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```

```
In [9]: # Convert categorical variables
insurance_df['smoker'] = insurance_df['smoker'].map({'yes': 1, 'no': 0})
insurance_df['sex'] = insurance_df['sex'].map({'male': 1, 'female': 0}) # A
insurance_df['region'] = insurance_df['region'].astype('category').cat.codes
```

```
In [13]: # Plot histograms for key variables
fig, axes = plt.subplots(2, 3, figsize=(15, 10))
fig.suptitle('Histograms of Key Variables')

columns = ['age', 'bmi', 'children', 'smoker', 'charges', 'region']
for i, col in enumerate(columns):
    sns.histplot(insurance_df[col], bins=30, kde=True, ax=axes[i//3, i%3])
    axes[i//3, i%3].set_title(f'Distribution of {col}')

plt.tight_layout()
plt.show()
```



```
In [15]: # Summary statistics
insurance_df.describe()
```

```
Out[15]:
```

	age	sex	bmi	children	smoker	region
count	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	0.505232	30.663397	1.094918	0.204783	1.515625
std	14.049960	0.500160	6.098187	1.205493	0.403694	1.104838
min	18.000000	0.000000	15.960000	0.000000	0.000000	0.000000
25%	27.000000	0.000000	26.296250	0.000000	0.000000	1.000000
50%	39.000000	1.000000	30.400000	1.000000	0.000000	2.000000
75%	51.000000	1.000000	34.693750	2.000000	0.000000	2.000000
max	64.000000	1.000000	53.130000	5.000000	1.000000	3.000000

```
In [ ]: #Look at mean, median, and standard deviation.
#Identify potential outliers in BMI and charges.
#check the skewness of charges (often right-skewed).
```

```
In [19]: # PMF for smokers vs. non-smokers (charges)
smoker_charges = insurance_df[insurance_df['smoker'] == 1]['charges']
non_smoker_charges = insurance_df[insurance_df['smoker'] == 0]['charges']


# Plot the PMF
plt.figure(figsize=(10, 5))
sns.histplot(smoker_charges, bins=30, kde=True, color='red', label='Smokers')
```

```

sns.histplot(non_smoker_charges, bins=30, kde=True, color='blue', label='Non-Smokers')

plt.legend()
plt.title("PMF of Insurance Charges for Smokers vs. Non-Smokers")
plt.xlabel("Charges")
plt.ylabel("Probability Density")
plt.show()

```



```

In [ ]: #Smokers have significantly higher insurance charges.
        #The distribution shifts right, indicating a strong effect.


```

```

In [23]: # Compute CDF
def plot_cdf(data, title):
    sorted_data = np.sort(data)
    y = np.arange(1, len(sorted_data) + 1) / len(sorted_data)
    plt.plot(sorted_data, y, marker='.', linestyle='none')
    plt.xlabel(title)
    plt.ylabel('CDF')
    plt.grid()

plt.figure(figsize=(8, 5))
plot_cdf(insurance_df['charges'], "CDF of Insurance Charges")
plt.show()

```




```
In [ ]: #A steep rise indicates where most people's charges are concentrated.
        #This helps understand the distribution and percentiles.
```

```
In [29]: # Fit a normal distribution to charges
mu, sigma = stats.norm.fit(insurance_df['charges'])

# Plot the histogram with the normal distribution curve
plt.figure(figsize=(10, 5))
sns.histplot(insurance_df['charges'], bins=30, kde=False, stat="density")
x = np.linspace(min(insurance_df['charges']), max(insurance_df['charges']), 100)
plt.plot(x, stats.norm.pdf(x, mu, sigma), label='Normal Distribution Fit', c='red')

plt.title('Checking Normality of Charges')
plt.xlabel('Charges')
plt.ylabel('Density')
plt.legend()
plt.show()
```





```
In [ ]: #The actual distribution is right-skewed, not normally distributed.
        #This confirms outliers affect charges.
```

```
In [33]: # Scatter plot: BMI vs. Charges
plt.figure(figsize=(8, 5))
sns.scatterplot(x=insurance_df['bmi'], y=insurance_df['charges'], hue=insurance_df['smoker'])
plt.title("BMI vs Charges")
plt.show()

# Scatter plot: Age vs. Charges
plt.figure(figsize=(8, 5))
sns.scatterplot(x=insurance_df['age'], y=insurance_df['charges'], hue=insurance_df['smoker'])
plt.title("Age vs Charges")
plt.show()

# Pearson correlation
print(insurance_df[['age', 'bmi', 'charges']].corr())
```



	age	bmi	charges
age	1.000000	0.109272	0.299008
bmi	0.109272	1.000000	0.198341
charges	0.299008	0.198341	1.000000

In []: *#BMI vs. Charges: A weak trend, but smokers have higher charges.
#Age vs. Charges: Slight positive correlation.*

```
In [37]: # Perform a t-test
#Testing if smokers have significantly higher charges.

t_stat, p_val = stats.ttest_ind(smoker_charges, non_smoker_charges)

print(f"T-Statistic: {t_stat}, P-Value: {p_val}")

# Interpretation
if p_val < 0.05:
    print("Reject Null Hypothesis: Smoking significantly affects charges.")
else:
    print("Fail to Reject Null Hypothesis: No significant effect.")
```

T-Statistic: 46.66492117272371, P-Value: 8.271435842179101e-283
Reject Null Hypothesis: Smoking significantly affects charges.

```
In [41]: # Simple Linear Regression
model = smf.ols('charges ~ bmi', data=insurance_df).fit()
print(model.summary())
```


OLS Regression Results

```

=====
==
Dep. Variable:          charges    R-squared:                0.0
39
Model:                  OLS      Adj. R-squared:            0.0
39
Method:                 Least Squares    F-statistic:            54.
71
Date:                   Sat, 01 Mar 2025    Prob (F-statistic):      2.46e-
13
Time:                   16:00:19    Log-Likelihood:         -1445
1.
No. Observations:      1338    AIC:                    2.891e+
04
Df Residuals:          1336    BIC:                    2.892e+
04
Df Model:               1
Covariance Type:       nonrobust
=====

```

```

=====
==
              coef      std err          t      P>|t|      [0.025      0.97
5]
-----
--
Intercept    1192.9372    1664.802      0.717      0.474    -2072.974    4458.8
49
bmi           393.8730      53.251      7.397      0.000      289.409      498.3
37
=====

```

```

=====
==
Omnibus:          261.030    Durbin-Watson:           1.9
83
Prob(Omnibus):    0.000    Jarque-Bera (JB):        431.0
91
Skew:             1.297    Prob(JB):                2.45e-
94
Kurtosis:         4.004    Cond. No.                 16
0.
=====

```

```

=====
==

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

In [45]: # Multiple Regression
model = smf.ols('charges ~ age + bmi + smoker', data=insurance_df).fit()
print(model.summary())

```

OLS Regression Results

```

=====
==
Dep. Variable:          charges    R-squared:                0.7
47
Model:                  OLS      Adj. R-squared:            0.7
47
Method:                 Least Squares    F-statistic:            131
6.
Date:                  Sat, 01 Mar 2025    Prob (F-statistic):      0.
00
Time:                  16:00:52    Log-Likelihood:          -1355
7.
No. Observations:      1338    AIC:                    2.712e+
04
Df Residuals:          1334    BIC:                    2.714e+
04
Df Model:               3
Covariance Type:       nonrobust
=====

```

```

=====
==
               coef      std err          t      P>|t|      [0.025      0.97
5]
-----
--
Intercept  -1.168e+04    937.569    -12.454    0.000    -1.35e+04    -9837.5
61
age         259.5475     11.934     21.748    0.000     236.136     282.9
59
bmi         322.6151     27.487     11.737    0.000     268.692     376.5
38
smoker      2.382e+04    412.867     57.703    0.000     2.3e+04     2.46e+
04
=====

```

```

=====
==
Omnibus:            299.709    Durbin-Watson:            2.0
77
Prob(Omnibus):      0.000    Jarque-Bera (JB):          710.1
37
Skew:               1.213    Prob(JB):                  6.25e-1
55
Kurtosis:           5.618    Cond. No.                  28
9.
=====

```

```

==
Notes:
[1] Standard Errors assume that the covariance matrix of the errors is corre
ctly specified.

```

```

In [37]: import statsmodels.api as sm

# Define independent variables (X) and dependent variable (y)
X = insurance_df[['age', 'bmi', 'smoker', 'region_northwest', 'region_southe
y = insurance_df['charges']

```

```
# Add a constant for the regression intercept  
X = sm.add_constant(X)
```

```
In [39]: # Fit the regression model  
model = sm.OLS(y, X).fit()  
  
# Display regression results  
print(model.summary())
```

OLS Regression Results

```

=====
==
Dep. Variable:          charges    R-squared:                0.7
49
Model:                  OLS        Adj. R-squared:            0.7
48
Method:                 Least Squares    F-statistic:              66
0.8
Date:                   Sat, 01 Mar 2025    Prob (F-statistic):       0.
00
Time:                   15:34:21    Log-Likelihood:           -1355
4.
No. Observations:      1338    AIC:                      2.712e+
04
Df Residuals:          1331    BIC:                      2.716e+
04
Df Model:               6
Covariance Type:       nonrobust
=====

```

```

=====
=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
const      -1.16e+04    976.200     -11.884    0.000    -1.35e+04    -
9686.503
age         258.6365     11.930      21.680    0.000      235.233
282.040
bmi         340.0076      28.673      11.858    0.000      283.759
396.256
smoker      2.385e+04    413.508      57.683    0.000      2.3e+04
2.47e+04
region_northwest -303.5207    477.850      -0.635    0.525     -1240.943
633.901
region_southeast -1038.6326    480.486      -2.162    0.031     -1981.225
-96.040
region_southwest -915.9394    479.558      -1.910    0.056     -1856.712
24.833
=====

```

```

=====
==
Omnibus:              298.282    Durbin-Watson:           2.0
79
Prob(Omnibus):        0.000    Jarque-Bera (JB):        705.0
89
Skew:                 1.208    Prob(JB):                7.80e-1
54
Kurtosis:             5.609    Cond. No.                30
7.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [ ]: #Smoking is the strongest predictor of charges.  
#BMI has a mild effect but interacts with smoking.  
#Age slightly influences charges.
```